



Google Dopamine

New Framework for Flexible and Reproducible Reinforcement Learning Research

Introduction

Reinforcement learning is an important domain of machine learning. Which mimics the human level learning. RL has gain a lot of momentum over the past few years. More and more research has been done in RL results in the improvements in the reinforcement learning methodolgies and techniques. Google has announced new Reinforcement learning tensor flow based framework called Dopamine. Which aims to provide flexibility, stability, and reproducibility for new and experienced RL researchers alike. Inspired by one of the main components in reward-motivated behaviour in the brain and reflecting the strong historical connection between neuroscience and reinforcement learning research, this platform aims to enable the kind of speculative research that can drive radical discoveries.

- Ease of use
- Reproducibility
- Benchmarking

Here we will code our simple agent using google Dopamine.

```
# @title Install necessary packages.
#dopamine for RL
!pip install --upgrade --no-cache-dir dopar
# dopamine dependencies
!pip install cmake
#Arcade Learning Environment
!pip install atari_py
```

Install necessary packages.



Collecting dopamine-rl

Downloading <https://files.pythonhosted.org/packages/a3/60/ce40162119275f8961b79ee1>

100% |██| 71kB 4.5MB/s

Requirement already satisfied, skipping upgrade: opencv-python>=3.4.1.15 in /usr/loc

Collecting gin-config>=0.1.1 (from dopamine-rl)

Downloading <https://files.pythonhosted.org/packages/e4/07/c8054ce483f058cd8fa2368c>

Requirement already satisfied, skipping upgrade: tensorflow in /usr/local/lib/python

Requirement already satisfied, skipping upgrade: absl-py>=0.2.2 in /usr/local/lib/py

Collecting gym>=0.10.5 (from dopamine-rl)

Downloading <https://files.pythonhosted.org/packages/9b/50/ed4a03d2be47ffd043be2ee5>

100% |██| 1.5MB 16.3MB/s

Requirement already satisfied, skipping upgrade: numpy>=1.11.3 in /usr/local/lib/pyt

Requirement already satisfied, skipping upgrade: six>=1.10.0 in /usr/local/lib/pytho

Requirement already satisfied, skipping upgrade: gast>=0.2.0 in /usr/local/lib/pytho

Requirement already satisfied, skipping upgrade: wheel>=0.26 in /usr/local/lib/pytho

Requirement already satisfied, skipping upgrade: grpcio>=1.8.6 in /usr/local/lib/pyt

Requirement already satisfied, skipping upgrade: protobuf>=3.6.0 in /usr/local/lib/p

Requirement already satisfied, skipping upgrade: termcolor>=1.1.0 in /usr/local/lib/

Requirement already satisfied, skipping upgrade: tensorboard<1.11.0,>=1.10.0 in /usr

Requirement already satisfied, skipping upgrade: astor>=0.6.0 in /usr/local/lib/pyth

Requirement already satisfied, skipping upgrade: setuptools<=39.1.0 in /usr/local/li

Requirement already satisfied, skipping upgrade: requests>=2.0 in /usr/local/lib/pyt

Collecting pygame>=1.2.0 (from gym>=0.10.5->dopamine-rl)

Downloading <https://files.pythonhosted.org/packages/1c/fc/dad5eaaab68f0c21e2f906a9>

100% |██| 1.0MB 26.8MB/s

Requirement already satisfied, skipping upgrade: markdown>=2.6.8 in /usr/local/lib/p

Requirement already satisfied, skipping upgrade: werkzeug>=0.11.10 in /usr/local/lib

Requirement already satisfied, skipping upgrade: urllib3<1.23,>=1.21.1 in /usr/local

Requirement already satisfied, skipping upgrade: certifi>=2017.4.17 in /usr/local/li

Requirement already satisfied, skipping upgrade: chardet<3.1.0,>=3.0.2 in /usr/local

Requirement already satisfied, skipping upgrade: idna<2.7,>=2.5 in /usr/local/lib/py

Requirement already satisfied, skipping upgrade: future in /usr/local/lib/python3.6/

Installing collected packages: gin-config, pygame, gym, dopamine-rl

Running setup.py install for gym ... done

```
# @title Necessary imports and globals
import numpy as np
import os
#DQN for baselines
from dopamine.agents.dqn import dqn_agent
from dopamine.atari import run_experiment
from dopamine.colab import utils as colab_utils
#warnings
from absl import flags
```

```
#where to store training logs
BASE_PATH = '/tmp/colab_dope_run' # @param
#which arcade environment?
GAME = 'Pong' # @param
```

Stored in directorv: /root/.cache/pdp/wheels/ac/79/85/b21b404d3469c3028aea3b7a1dte

```
# @title Create a new agent from scratch.
```

```
#define where to store log data
LOG_PATH = os.path.join(BASE_PATH, 'basic_;
```

```
class BasicAgent(object):
    """This agent randomly selects an action
    actions with probability switch_prob."""
    def __init__(self, sess, num_actions, sw:
        #tensorflow session
```

Necessary imports and globals

BASE_PATH: '/tmp/colab_dope_run'

GAME: 'Pong'

Create a new agent from scratch.

```

self._sess = sess
#how many possible actions can it take
self._num_actions = num_actions
# probability of switching actions in 1
self._switch_prob = switch_prob
#initialize the action to take (random)
self._last_action = np.random.randint(0, num_actions)
#not debugging
self.eval_mode = False

#How select an action?
#we define our policy here
def _choose_action(self):
    if np.random.random() <= self._switch_prob:
        self._last_action = np.random.randint(0, num_actions)
    return self._last_action

#when it checkpoints during training, any unused data is discarded
def bundle_and_checkpoint(self, unused_checkpoint_dir, unused_data):
    pass

#loading from checkpoint
def unbundle(self, unused_checkpoint_dir, unused_data):
    pass

#first action to take
def begin_episode(self, unused_observation):
    return self._choose_action()

#cleanup
def end_episode(self, unused_reward):
    pass

#we can update our policy here
#using the reward and observation
#dynamic programming, Q learning, monte carlo
def step(self, reward, observation):
    return self._choose_action()

def create_basic_agent(sess, environment):
    """The Runner class will expect a function that returns a tuple of (observation, reward, done, info)"""
    return BasicAgent(sess, num_actions=environment.get_num_actions(),
                       switch_prob=0.2)

# Create the runner class with this agent.
# to terminate quickly, as this is mostly for testing
# use the framework. We also explicitly terminate after 200 episodes (out of the standard 200) to demonstrate the framework
basic_runner = run_experiment.Runner(LOG_PATH, create_basic_agent,
                                     game_name='CartPole-v0',
                                     num_episodes=200,
                                     training_dir='.',
                                     eval_dir='.',
                                     max_episode_steps=1000)

# @title Train Basic Agent.
print('Will train basic agent, please be patient')
basic_runner.run_experiment()
print('Done training!')

```



Train Basic Agent.

```

Will train basic agent, please be patient, may be a while...
INFO:tensorflow:Beginning training...
INFO:tensorflow:Starting iteration 0
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 676.90
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 1
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 697.19
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 2
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 704.06
INFO:tensorflow:Average undiscounted return per evaluation episode: -1.00
INFO:tensorflow:Starting iteration 3
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 695.77
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 4
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 690.75
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 5
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 700.27
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 6
INFO:tensorflow:Average undiscounted return per training episode: -1.00
INFO:tensorflow:Average training steps per second: 699.86
INFO:tensorflow:Average undiscounted return per evaluation episode: -1.00
INFO:tensorflow:Starting iteration 7
INFO:tensorflow:Average undiscounted return per training episode: -1.00
INFO:tensorflow:Average training steps per second: 673.64
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 8
INFO:tensorflow:Average undiscounted return per training episode: -1.00
INFO:tensorflow:Average training steps per second: 663.00
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 9
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 692.02
INFO:tensorflow:Average undiscounted return per evaluation episode: -1.00
INFO:tensorflow:Starting iteration 10
INFO:tensorflow:Average undiscounted return per training episode: -2.00
INFO:tensorflow:Average training steps per second: 685.70
INFO:tensorflow:Average undiscounted return per evaluation episode: -1.00
INFO:tensorflow:Starting iteration 11
INFO:tensorflow:Average undiscounted return per training episode: 0.00
INFO:tensorflow:Average training steps per second: 683.18
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00
INFO:tensorflow:Starting iteration 12
INFO:tensorflow:Average undiscounted return per training episode: -1.00
INFO:tensorflow:Average training steps per second: 692.66
INFO:tensorflow:Average undiscounted return per evaluation episode: -2.00

```

```
# @title Load baseline data# @titl
```

```
!gsutil -q -m cp -R gs://download-dopamine-rl/preprocessed-benchmarks/* /content/
```

```
experimental_data = colab_utils.load_baselines('/content')
```

```
# @title Load the training logs.
```

Load the training logs.

```
basic_data = colab_utils.read_experiment(1)
basic_data['agent'] = 'BasicAgent'
basic_data['run_number'] = 1
experimental_data[GAME] = experimental_data[GAME].load(basic_data)
```


➞ Reading statistics from: /tmp/colab_dope_run/basic_agent/Pong//logs/log_199

```
# @title Plot training results.
```

Plot training results.

```
import seaborn as sns
import matplotlib.pyplot as plt

fig, ax = plt.subplots(figsize=(16,8))
sns.tsplot(data=experimental_data[GAME], t:
            condition='agent', value='train_
plt.title(GAME)
plt.show()
```



So is the training results of agent which we trained above.

You can find this in github too.
Follow AI based punlication on Medium ****"The 21st Century"**
<https://medium.com/the-21st-century>

****Hamza Abdullah ****

You can find this in github too. Follow AI based
punlication on Medium **"The 21st Century"**
<https://medium.com/the-21st-century>

Hamza Abdullah